

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")
%matplotlib inline
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Ridge, Lasso,
ElasticNet
from sklearn.metrics import r2_score, mean_absolute_error,
mean_squared_error
from scipy import stats
from statsmodels.stats.outliers_influence import
variance_inflation_factor
import statsmodels.api as sm
import statsmodels.stats.api as sms

jb=pd.read_csv("C:/Users/shrad/Downloads/Jamboree_Admission.csv")
jb.head(10)

```

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR
CGPA \						
0	1	337	118	4	4.5	4.5
9.65						
1	2	324	107	4	4.0	4.5
8.87						
2	3	316	104	3	3.0	3.5
8.00						
3	4	322	110	3	3.5	2.5
8.67						
4	5	314	103	2	2.0	3.0
8.21						
5	6	330	115	5	4.5	3.0
9.34						
6	7	321	109	3	3.0	4.0
8.20						
7	8	308	101	2	3.0	4.0
7.90						
8	9	302	102	1	2.0	1.5
8.00						
9	10	323	108	3	3.5	3.0
8.60						

	Research	Chance of Admit
0	1	0.92
1	1	0.76
2	1	0.72
3	1	0.80

4	0	0.65
5	1	0.90
6	1	0.75
7	0	0.68
8	0	0.50
9	0	0.45

```
jb.drop(columns="Serial No.",inplace=True)
```

Insight We can remove the unnecessary column "Serial No." from data, as it doesn't have any contribution for data visualization and operation.

```
jb.isnull().sum()
```

GRE Score	0
TOEFL Score	0
University Rating	0
SOP	0
LOR	0
CGPA	0
Research	0
Chance of Admit	0

dtype: int64

Insight There are null values in our data.

```
jb.nunique()
```

GRE Score	49
TOEFL Score	29
University Rating	5
SOP	9
LOR	9
CGPA	184
Research	2
Chance of Admit	61

dtype: int64

Insight Features like "Research","University Rating","SOP","LOR" are categorical features.

```
jb.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   GRE Score              500 non-null    int64
1   TOEFL Score            500 non-null    int64
```

2	University Rating	500	non-null	int64
3	SOP	500	non-null	float64
4	LOR	500	non-null	float64
5	CGPA	500	non-null	float64
6	Research	500	non-null	int64
7	Chance of Admit	500	non-null	float64

dtypes: float64(4), int64(4)

memory usage: 31.4 KB

jb.describe()

	GRE Score	TOEFL Score	University Rating	SOP
count	500.000000	500.000000	500.000000	500.000000
mean	316.472000	107.192000	3.114000	3.374000
std	11.295148	6.081868	1.143512	0.991004
min	290.000000	92.000000	1.000000	1.000000
25%	308.000000	103.000000	2.000000	2.500000
50%	317.000000	107.000000	3.000000	3.500000
75%	325.000000	112.000000	4.000000	4.000000
max	340.000000	120.000000	5.000000	5.000000

	CGPA	Research	Chance of Admit
count	500.000000	500.000000	500.000000
mean	8.576440	0.560000	0.72174
std	0.604813	0.496884	0.14114
min	6.800000	0.000000	0.34000
25%	8.127500	0.000000	0.63000
50%	8.560000	1.000000	0.72000
75%	9.040000	1.000000	0.82000
max	9.920000	1.000000	0.97000

mean=jb.mean(axis=0)

median=jb.median(axis=0)

mode=pd.DataFrame(jb.mode(axis=0), columns=jb.columns)

print(mean, '\n----- mean:-----\n', median, '\n-----

mode:-----\n')

mode.head()

std=jb.describe().loc['std']

Q1=jb.describe().loc['25%']

Q3=jb.describe().loc['75%']

IQR=Q3-Q1

```

max=Q1=jb.describe().loc['max']
min=Q1=jb.describe().loc['min']
print(std, '\n----- IQR:-----\n', IQR, '\n----- max:\n', max, '\n***** min:*****\n'), min
skew=jb.skew(axis=1)
kurt=pd.DataFrame(jb.kurt(axis=1), columns=["kurt"])
kurt.head()
#print('\n***** skew:*****\n', skew, '\n***** kurt:*****\n', kurt)

```

```

GRE Score      316.47200
TOEFL Score    107.19200
University Rating  3.11400
SOP            3.37400
LOR            3.48400
CGPA           8.57644
Research       0.56000
Chance of Admit  0.72174
dtype: float64

```

```

----- mean:-----
GRE Score      317.00
TOEFL Score    107.00
University Rating  3.00
SOP            3.50
LOR            3.50
CGPA           8.56
Research       1.00
Chance of Admit  0.72
dtype: float64

```

```

----- mode:-----
GRE Score      11.295148
TOEFL Score    6.081868
University Rating  1.143512
SOP            0.991004
LOR            0.925450
CGPA           0.604813
Research       0.496884
Chance of Admit  0.141140
Name: std, dtype: float64

```

```

----- IQR:-----
GRE Score      17.0000
TOEFL Score    9.0000
University Rating  2.0000
SOP            1.5000
LOR            1.0000
CGPA           0.9125
Research       1.0000
Chance of Admit  0.1900
dtype: float64

```

```

----- max:
GRE Score          340.00
TOEFL Score        120.00
University Rating   5.00
SOP                5.00
LOR                5.00
CGPA               9.92
Research           1.00
Chance of Admit    0.97
Name: max, dtype: float64
***** min:*****

```

```

      kurt
0  5.342020
1  5.607376
2  5.608048
3  5.439713
4  5.605129

```

```
jb.cov()
```

	GRE Score	TOEFL Score	University Rating
SOP \			
GRE Score	127.580377	56.825026	8.206605
6.867206			
TOEFL Score	56.825026	36.989114	4.519150
3.883960			
University Rating	8.206605	4.519150	1.307619
0.825014			
SOP	6.867206	3.883960	0.825014
0.982088			
LOR	5.484521	3.048168	0.644112
0.608701			
CGPA	5.641944	2.981607	0.487761
0.426845			
Research	3.162004	1.411303	0.242645
0.200962			
Chance of Admit	1.291862	0.680046	0.111384
0.095691			

	LOR	CGPA	Research	Chance of Admit
GRE Score	5.484521	5.641944	3.162004	1.291862
TOEFL Score	3.048168	2.981607	1.411303	0.680046
University Rating	0.644112	0.487761	0.242645	0.111384
SOP	0.608701	0.426845	0.200962	0.095691
LOR	0.856457	0.356807	0.171303	0.084296
CGPA	0.356807	0.365799	0.150655	0.075326
Research	0.171303	0.150655	0.246894	0.038282
Chance of Admit	0.084296	0.075326	0.038282	0.019921



```
jb.corr()
```

	GRE Score	TOEFL Score	University Rating	SOP
\				
GRE Score	1.000000	0.827200	0.635376	0.613498
TOEFL Score	0.827200	1.000000	0.649799	0.644410
University Rating	0.635376	0.649799	1.000000	0.728024
SOP	0.613498	0.644410	0.728024	1.000000
LOR	0.524679	0.541563	0.608651	0.663707
CGPA	0.825878	0.810574	0.705254	0.712154
Research	0.563398	0.467012	0.427047	0.408116
Chance of Admit	0.810351	0.792228	0.690132	0.684137

	LOR	CGPA	Research	Chance of Admit
GRE Score	0.524679	0.825878	0.563398	0.810351
TOEFL Score	0.541563	0.810574	0.467012	0.792228
University Rating	0.608651	0.705254	0.427047	0.690132
SOP	0.663707	0.712154	0.408116	0.684137
LOR	1.000000	0.637469	0.372526	0.645365
CGPA	0.637469	1.000000	0.501311	0.882413
Research	0.372526	0.501311	1.000000	0.545871
Chance of Admit	0.645365	0.882413	0.545871	1.000000


Univariate Analysis

```
plt.figure(figsize=(20,8))
plt.subplot(1,2,1)
researchcount=jb["Research"].value_counts()
ax=sns.countplot(data=jb,x="Research")
for bars in ax.containers:
    ax.bar_label(bars)
plt.subplot(1,2,2)
plt.pie(
researchcount,
labels=researchcount.index,
startangle=90,
autopct="%.2f")
plt.show()
```



Insight More than half of the have research which is 56%.

```
plt.figure(figsize=(20,8))
plt.subplot(1,2,1)
counts, bins, patches = plt.hist(jb['TOEFL Score'], bins=6, alpha=0.7,
color='green', rwidth=0.8)
# Calculate and annotate percentages
total = len(jb)
```



```
plt.figure(figsize=(12, 6))

# 1. Create a histogram for TOEFL scores
plt.subplot(1, 2, 1)
counts, bins, patches = plt.hist(jb['TOEFL Score'], bins=10,
color='blue', alpha=0.7)
total = sum(counts)

# Add percentage labels on top of the bars
for count, patch in zip(counts, patches):
    height = patch.get_height()
    plt.text(patch.get_x() + patch.get_width() / 2, height + 1,
             '{:.1f}%'.format(100 * count / total), ha='center')

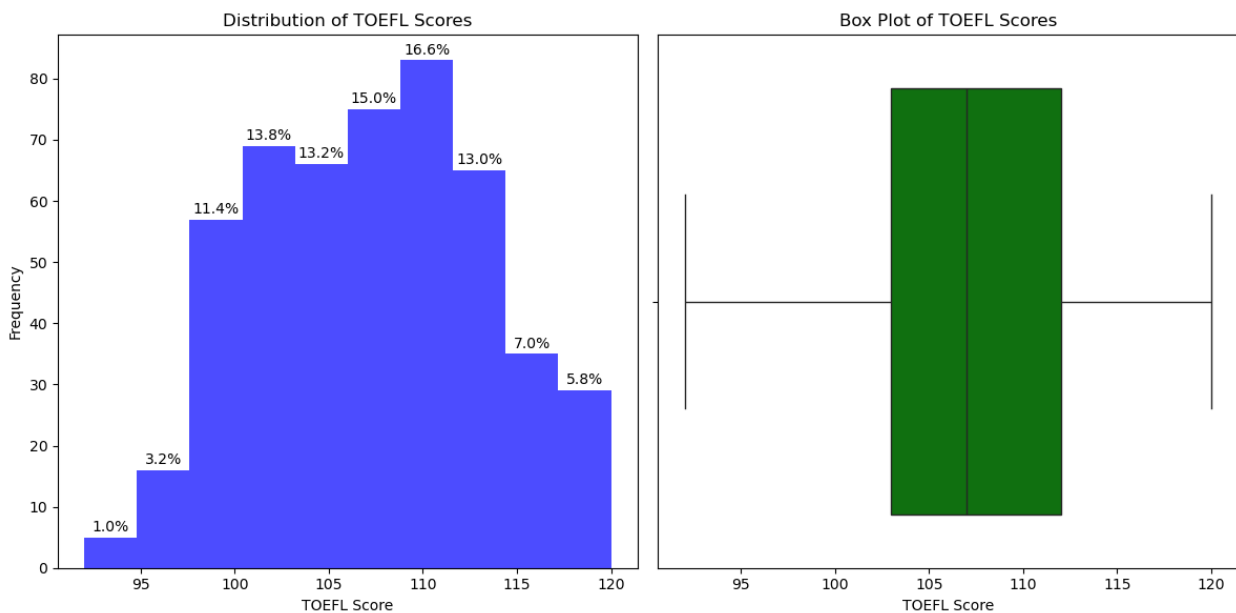
# Set labels and title for the histogram
plt.xlabel('TOEFL Score')
plt.ylabel('Frequency')
```



```
plt.title('Distribution of TOEFL Scores')

# 2. Create a box plot for TOEFL scores
plt.subplot(1, 2, 2)
sns.boxplot(data=jb, x="TOEFL Score", color="green")
plt.title('Box Plot of TOEFL Scores')



# Show the plots
plt.tight_layout()
plt.show()
```



Insight

1. About 50% students score about 100-110.
2. There is no outliers in TOEFL Score.



```
plt.figure(figsize=(20,8))
plt.subplot(1,2,1)
#jb['GRE Score'].plot(kind='hist', bins=6, alpha=0.7, color='green',
#figsize=(10, 6), width=2.5)
counts, bins, patches = plt.hist(jb['GRE Score'], bins=6, alpha=0.7,
color='green', rwidth=0.8)
# Calculate and annotate percentages
total = len(jb)
for count, patch in zip(counts, patches):
    height = patch.get_height()
    plt.text(patch.get_x() + patch.get_width() / 2, height + 1, '{:.1f}'
%'.format(100 * count / total), ha='center')
plt.subplot(1,2,2)
sns.boxplot(data=jb,x="GRE Score",color="green")
plt.show()
```



Insight

1. About 49% of students scored 308-322 in GRE.
2. 9.4% students scored 332-339 and 6.2% scored 291-228.
3. There are no outliers in GRE Score.


```
plt.figure(figsize=(20,8))
plt.subplot(1,2,1)
#jb['SOP'].plot(kind='hist', bins=9, alpha=0.7, color='green',
figsize=(10, 6),width=0.4)
counts, bins, patches = plt.hist(jb['SOP'], bins=6, alpha=0.7,
color='green', rwidth=0.8)
# Calculate and annotate percentages
total = len(jb)
for count, patch in zip(counts, patches):
    height = patch.get_height()
    plt.text(patch.get_x() + patch.get_width() / 2, height + 1, '{:.1f}%'
.format(100 * count / total), ha='center')
plt.subplot(1,2,2)
sns.boxplot(data=jb,x="SOP",color="green")
plt.show()
```




Insight

1. About 50% of students have SOP score around 3.5 to 4.0.
2. There are no outliers in GRE Score.

```
researchcount=jb["University Rating"].value_counts()
ax1=sns.countplot(data=jb,x="University Rating")
for bars in ax1.containers:
    ax1.bar_label(bars)
```



```
jb.columns = jb.columns.str.strip()
sns.pairplot(data=jb, y_vars=["Chance of Admit"])
plt.show()
```




Insight

1. By the above series of graph, chance of admit increases by increment in GRE Score. Same trend can be seen in TOEFL Score and CGPA.
2. On the other hand, University Rating, SOP, LOR, Research doesn't show any trend in chance of admit.

Bivariate Analysis

Correlation Analysis

```
plt.figure(figsize=(16,8))
sns.heatmap(jb.corr(), annot=True, cmap='Reds', linewidths=0.1)
plt.show()
```



Insight

1. CGPA have the highest correlation with the chance of admission.
2. Research have the lowest correlation with the chance of admission.

```

scaler = StandardScaler()
scaled_jb = pd.DataFrame(scaler.fit_transform(jb), columns =
jb.columns)
scaled_jb

```

	GRE Score	TOEFL Score	University Rating	SOP	LOR
CGPA \					
0	1.819238	1.778865	0.775582	1.137360	1.098944
1	0.667148	-0.031601	0.775582	0.632315	1.098944
2	-0.041830	-0.525364	-0.099793	-0.377773	0.017306
3	0.489904	0.462163	-0.099793	0.127271	-1.064332
4	-0.219074	-0.689952	-0.975168	-1.387862	-0.523513
...
495	1.376126	0.132987	1.650957	1.137360	0.558125
496	1.819238	1.614278	1.650957	1.642404	1.639763
497	1.198882	2.108041	1.650957	1.137360	1.639763
498	-0.396319	-0.689952	0.775582	0.632315	1.639763
499	0.933015	0.955926	0.775582	1.137360	1.098944
Research		Chance of Admit			
0	0.886405	1.406107			
1	0.886405	0.271349			
2	0.886405	-0.012340			
3	0.886405	0.555039			
4	-1.128152	-0.508797			
...			
495	0.886405	1.051495			
496	0.886405	1.689797			
497	0.886405	1.477030			
498	-1.128152	0.058582			
499	-1.128152	0.838728			

```

[500 rows x 8 columns]

```

Splitting data for training and testing

```

x=scaled_jb.iloc[:, :-1]
y=scaled_jb.iloc[:, -1]
print(x.shape,y.shape)

(500, 7) (500,)

x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=0.2,random_state=42)
print(f'Shape of x_train: {x_train.shape}')
print(f'Shape of x_test: {x_test.shape}')
print(f'Shape of y_train: {y_train.shape}')
print(f'Shape of y_test: {y_test.shape}')

Shape of x_train: (400, 7)
Shape of x_test: (100, 7)
Shape of y_train: (400,)
Shape of y_test: (100,)

```

Linear Regression

```

lr_model = LinearRegression()
lr_model.fit(x_train,y_train)

LinearRegression()

y_pred_train = lr_model.predict(x_train)
y_pred_test = lr_model.predict(x_test)

```

R2 score on train data

```

r2=r2_score(y_train,y_pred_train)
print("r2 score-> ",r2)
lr=lr_model.score(x_train,y_train)
print("lr score-> ",lr)

r2 score-> 0.8210671369321554
lr score-> 0.8210671369321554

```

R2 score on test data

```

r2_score(y_test,y_pred_test)
print("r2 score-> ",r2)
lr=lr_model.score(x_test,y_test)
print("lr score-> ",lr)

r2 score-> 0.8210671369321554
lr score-> 0.8188432567829628

```

All features coefficients and features

```
lr_model_weights = pd.DataFrame(lr_model.coef_.reshape(1, -
1), columns=jb.columns[:-1])
lr_model_weights["Intercept"] = lr_model.intercept_
lr_model_weights
```

	GRE Score	TOEFL Score	University Rating	SOP	LOR
CGPA \					
0	0.194823	0.129095	0.020812	0.012735	0.113028
0.482199					

	Research	Intercept
0	0.084586	0.007736

Insight

1. CGPA,GRE Score,TOEFL Score have highest weights.
2. University Rating, SOP, Research have lowest weights.
3. Intercept (w0) is very low

```
import numpy as np
from sklearn.metrics import mean_squared_error, mean_absolute_error,
r2_score

def model_evaluation(y_actual, y_forecast, model):
    n = len(y_actual)

    # Determine the number of coefficients
    if len(model.coef_.shape) == 1:
        p = len(model.coef_)
    else:
        p = len(model.coef_[0])

    # Calculate error metrics
    MSE = np.round(mean_squared_error(y_true=y_actual,
y_pred=y_forecast, squared=True), 2)
    MAE = np.round(mean_absolute_error(y_true=y_actual,
y_pred=y_forecast), 2)
    RMSE = np.round(mean_squared_error(y_true=y_actual,
y_pred=y_forecast, squared=False), 2)

    # Calculate R2 and Adjusted R2
    r2 = np.round(r2_score(y_true=y_actual, y_pred=y_forecast), 2)
    adj_r2 = np.round(1 - ((1 - r2) * (n - 1) / (n - p - 1)), 2)

    # Print the evaluation metrics
    print(f"MSE: {MSE}\nMAE: {MAE}\nRMSE: {RMSE}\nR2 Score: {r2}\
nAdjusted R2: {adj_r2}")

# Example usage (you would replace y_actual, y_forecast, and model
```

```

with your actual data)
# model_evaluation(y_actual, y_forecast, model)

model_evaluation(y_train.values, y_pred_train, lr_model)

MSE: 0.18
MAE: 0.3
RMSE: 0.42
R2 Score: 0.82
Adjusted R2: 0.82

model_evaluation(y_test.values, y_pred_test, lr_model)

MSE: 0.19
MAE: 0.3
RMSE: 0.43
R2 Score: 0.82
Adjusted R2: 0.81

```

Linear Regression using OLS

```

new_x_train = sm.add_constant(x_train)
model = sm.OLS(y_train, new_x_train)
results = model.fit()
# statistical summary of the model
print(results.summary())

```

OLS Regression Results

```

=====
=====
Dep. Variable:          Chance of Admit    R-squared:
0.821
Model:                                OLS    Adj. R-squared:
0.818
Method:                Least Squares    F-statistic:
257.0
Date:                  Mon, 28 Oct 2024    Prob (F-statistic):
3.41e-142
Time:                  10:55:50    Log-Likelihood:
-221.69
No. Observations:      400    AIC:
459.4
Df Residuals:          392    BIC:
491.3
Df Model:              7

Covariance Type:      nonrobust

=====
=====

```


[0.025		0.975]	coef	std err	t	P> t	

const			0.0077	0.021	0.363	0.717	-
0.034		0.050					
GRE Score			0.1948	0.046	4.196	0.000	
0.104		0.286					
TOEFL Score			0.1291	0.041	3.174	0.002	
0.049		0.209					
University Rating			0.0208	0.034	0.611	0.541	-
0.046		0.088					
SOP			0.0127	0.036	0.357	0.721	-
0.057		0.083					
LOR			0.1130	0.030	3.761	0.000	
0.054		0.172					
CGPA			0.4822	0.046	10.444	0.000	
0.391		0.573					
Research			0.0846	0.026	3.231	0.001	
0.033		0.136					
=====							
=====							
Omnibus:			86.232	Durbin-Watson:			
2.050							
Prob(Omnibus):			0.000	Jarque-Bera (JB):			
190.099							
Skew:			-1.107	Prob(JB):			
5.25e-42							
Kurtosis:			5.551	Cond. No.			
5.72							
=====							
=====							
Notes:							
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.							

Testing Assumptions of Linear Regression Model

1. No multicollinearity: Multicollinearity check by VIF(Variance Inflation Factor) score. Variables are dropped one-by-one till none has a VIF>5.
2. Mean of Residuals should be close to zero.
3. Linear relationship between independent & dependent variables. This can be checked using the following methods: Scatter plots Regression plots Pearson Correlation
4. Test for Homoscedasticity Create a scatterplot of residuals against predicted values. `r2 = np.round(r2_score(y_true=y_actual, y_pred=y_forecast),2)` `adj_r2 = np.round(1 - ((1-r2)*(n-1)/(n-p-1)),2)` `return print(f"MSE: {MSE}\nMAE: {MAE}\nRMSE: {RMSE}\nR2 Score: {r2}\nAdjusted R2: {adj_r2}")` `model_evaluation(y_train.values, y_pred_train, lr_model)`

```
model_evaluation(y_test.values, y_pred_test, lr_model)
```

```
new_x_train = sm.add_constant(x_train) model = sm.OLS(y_train, new_x_train) results = model.fit()
```

statstical summary of the model

```
print(results.summary())
```

Perform a Goldfeld-Quandt test to check the presence of Heteroscedasticity in the data.

- If the obtained **p-value** > 0.05, there is no strong evidence of heteroscedasticity.
1. Normality of Residuals Almost bell-shaped curve in residuals distribution.
 2. Impact of Outliers

Multicollinearity check:

VIF (Variance Inflation Factor) is a measure that quantifies the severity of multicollinearity in a regression analysis. It assesses how much the variance of the estimated regression coefficient is inflated due to collinearity. The formula for VIF is as follows: $VIF(j) = 1 / (1 - R(j)^2)$

Where:

j represents the jth predictor variable.

$R(j)^2$ is the coefficient of determination (R-squared) obtained from regressing the jth predictor variable on all the other predictor variables. "

Calculate the VIF for each variable. Identify variables with VIF greater than 5. Drop the variable with the highest VIF. Repeat steps 1-3 until no variable has a VIF greater than 5.

```
vif = pd.DataFrame()
vif['Variable'] = x_train.columns
vif['VIF'] = [variance_inflation_factor(x_train.values, i) for i in range(x_train.shape[1])]
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

	Variable	VIF
5	CGPA	4.653698
0	GRE Score	4.489201
1	TOEFL Score	3.665067
3	SOP	2.785753
2	University Rating	2.571847
4	LOR	1.977668
6	Research	1.517206


Insight

As the Variance Inflation Factor(VIF) score is less than 5 for all the features we can say that there is no much multicollinearity between the features.

Mean Of Residuals


1. If mean of residuals is significantly non-zero, then the model is overestimating or underestimating the observed values.
2. If the mean of residuals is close to zero then on average predictions made by linear regression model are accurate, within the equal balance of overestimating and underestimating. This is the desired characteristics for well-fitted regression model.

```
residual = y_test.values - y_pred_test  
  
residual_train = y_train.values - y_pred_train  
residual_train.mean()  
  
3.7192471324942746e-17  
  
residual.mean()  
  
-0.03867840379282768  
  
plt.figure(figsize=(15,5))  
sns.histplot(residual, kde= True,color='r')  
plt.title('Residuals Hist')  
sns.despine(left=True)  
plt.ylabel("")  
plt.yticks([])  
plt.show()
```



```
plt.figure(figsize=(15,5))  
sns.histplot(residual_train, kde= True,color='r')  
plt.title('Residuals Hist')  
sns.despine(left=True)
```

```
plt.ylabel("")
plt.yticks([])
plt.show()
```




Insight

The mean of residual is close to zero, therefore our model is unbiased

Linear Relationships:

```
plt.figure(figsize=(15,5))
plt.subplot(121)
plt.title('Residuals_train data')
sns.regplot(x=y_pred_train, y=residual_train, lowess=True,
color='r',line_kws={'color': 'green'})
plt.axhline(y=0, color='k', linestyle='--')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.subplot(122)
plt.title('Residuals_test data')
sns.regplot(x=y_pred_test, y=residual,
lowess=True,color='r',line_kws={'color': 'green'})
plt.axhline(y=0, color='k', linestyle='--')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
sns.despine()
plt.show()
```

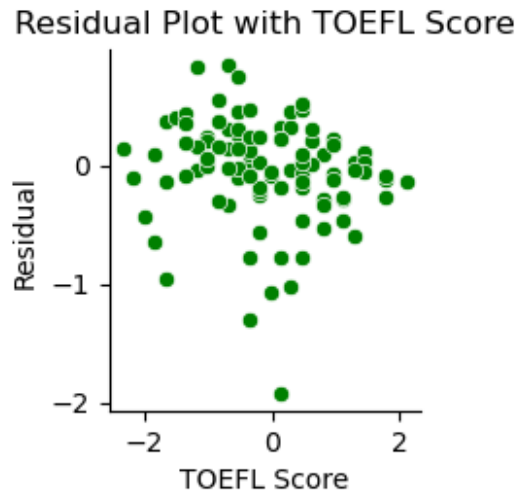


Insights:

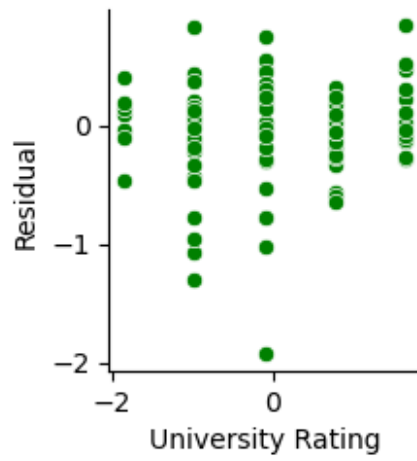
1. From the Joint plot & pairplot in the graphical analysis, we can say that there is linear relationship between dependent variable and independent variables.

Homoscedacity

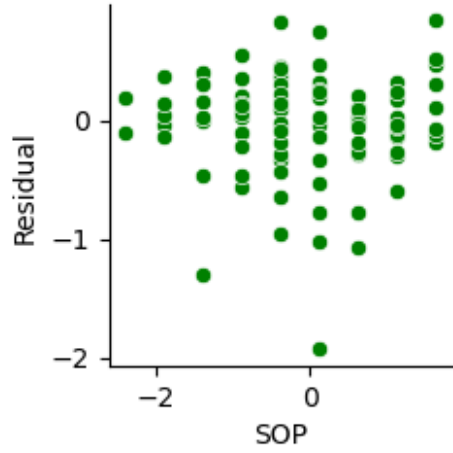
```
# Scatterplot of residuals with each independent variable to check for
Homoscedasticity
plt.figure(figsize=(15,8))
i=1
for col in x_test.columns[:-1]:
    plt.subplot(2,3,i)
    sns.scatterplot(x=x_test[col].values.reshape((-1,)),
y=residual.reshape((-1,)),color='g')
    plt.title(f'Residual Plot with {col}')
    plt.xlabel(col)
    plt.ylabel('Residual')
    i+=1
plt.tight_layout()
sns.despine()
plt.show();
```




Residual Plot with University Rating

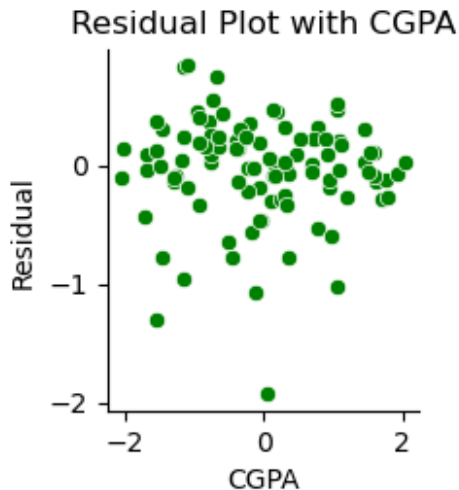


Residual Plot with SOP



Residual Plot with LOR





```
ols_model = results
predicted = ols_model.predict()
residuals = ols_model.resid
```

Breusch-Pagan test for Homoscedasticity

Null Hypothesis -- H_0 : Homoscedasticity is present in residuals. Alternate Hypothesis -- H_a : Heteroscedasticity is present in residuals. α : 0.05

Insights


Since we do not see any significant change in the spread of residuals with respect to change in independent variables, we can conclude that Homoscedasticity is met. Since the p-value is much lower than the alpha value, we can Reject the null hypothesis and conclude that *Heteroscedasticity is present*. Since the p-value is significantly less than the conventional significance level (e.g., 0.05), we reject the null hypothesis of homoscedasticity. This suggests that there is evidence of heteroscedasticity in the residuals, indicating that the variance of the residuals is not constant across all levels of the independent variables. This violation of the homoscedasticity assumption may affect the validity of the linear regression model's results.

Normality of Residuals:

To check normality, we will follow below methods:-


1. Residual Histogram
2. Q-Q Plot
3. Anderson-Darling or Jarque_Bera Test


```
plt.figure(figsize=(15,5))
sns.histplot(residual, kde=True,color='r')
plt.title('Residuals Hist')
sns.despine(left=True)
plt.ylabel("")
plt.yticks([])
plt.show()
```



```
plt.figure(figsize=(15,5))
sm.qqplot(residual,line='45')
plt.title('QQ Plot of Residuals')
plt.ylabel('Residual Quantiles')
sns.despine()
plt.show()
```

<Figure size 1500x500 with 0 Axes>



JARQUE BERA test:

```
jb_stat, jb_p_value = stats.jarque_bera(residual)
print("Jarque-Bera Test Statistic:", jb_stat)
print("p-value:", jb_p_value)
if jb_p_value < 0.05:
    print("Reject the null hypothesis: Residuals are not normally distributed.")
else:
    print("Fail to reject the null hypothesis: Residuals are normally distributed.")
```

Jarque-Bera Test Statistic: 74.10190609972128
p-value: 8.109153870348849e-17
Reject the null hypothesis: Residuals are not normally distributed.

Jarque-Bera Test Statistic: 74.10190609972094 p-value: 8.109153870350212e-17 Reject the null hypothesis: Residuals are not normally distributed.

Insight

1. From Hisplot and Kdeplot we can say that Residuals are left skewed.
2. The QQ plot shows that residuals are slightly deviating from the straight diagonal , thus not Gaussian.
3. From Jarque Bera test , we conclude that the Residuals are Not Normally distributed. Hence this assumption is not met.

Lasso and Ridge Regression - L1 & L2 Regularization

Lasso Regression:

```
model_lasso = Lasso(alpha=0.45)

model_lasso = Lasso(alpha=0.45)
model_lasso.fit(x_train, y_train)

Lasso(alpha=0.45)

model_ridge = Ridge()
model_ridge.fit(x_train, y_train)

Ridge()

y_pred_train_ridge = model_ridge.predict(x_train)
y_pred_test_ridge = model_ridge.predict(x_test)
y_pred_train_lasso = model_lasso.predict(x_train)
y_pred_test_lasso = model_lasso.predict(x_test)

lasso_model_weights = pd.DataFrame(model_lasso.coef_.reshape(1, -1), columns=jb.columns[:-1])
lasso_model_weights["Intercept"] = model_lasso.intercept_
lasso_model_weights
```

	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA
Research \						
0	0.019231	0.0	0.0	0.0	0.0	0.408647
0.0						
Intercept						
0	0.013919					

```
ridge_model_weights = pd.DataFrame(model_ridge.coef_.reshape(1, -1), columns=jb.columns[:-1])
ridge_model_weights["Intercept"] = model_ridge.intercept_
ridge_model_weights
```

	GRE Score	TOEFL Score	University Rating	SOP	LOR
CGPA \					
0	0.195584	0.130073	0.021575	0.013802	0.113221
	0.478123				

	Research	Intercept
0	0.084673	0.007726

```

print('Linear Regression Training Accuracy\n')
model_evaluation(y_train.values, y_pred_train, lr_model)
print('- '*25)
print('\nLinear Regression Test Accuracy\n')
model_evaluation(y_test.values, y_pred_test, lr_model)
print('--- '*25)
print('\nRidge Regression Training Accuracy\n')
model_evaluation(y_train.values, y_pred_train_ridge, model_ridge)
print('- '*25)
print('\n\nRidge Regression Test Accuracy\n')
model_evaluation(y_test.values, y_pred_test_ridge, model_ridge)
print('--- '*25)
print('\n\nLasso Regression Training Accuracy\n')
model_evaluation(y_train.values, y_pred_train_lasso, model_lasso)
print('- '*25)
print('\n\nLasso Regression Test Accuracy\n')
model_evaluation(y_test.values, y_pred_test_lasso, model_lasso)
print('--- '*25)

```

Linear Regression Training Accuracy

MSE: 0.18
 MAE: 0.3
 RMSE: 0.42
 R2 Score: 0.82
 Adjusted R2: 0.82

Linear Regression Test Accuracy

MSE: 0.19
 MAE: 0.3
 RMSE: 0.43
 R2 Score: 0.82
 Adjusted R2: 0.81

Ridge Regression Training Accuracy

MSE: 0.18
 MAE: 0.3

```
RMSE: 0.42
R2 Score: 0.82
Adjusted R2: 0.82
-----
```

Ridge Regression Test Accuracy

```
MSE: 0.19
MAE: 0.3
RMSE: 0.43
R2 Score: 0.82
Adjusted R2: 0.81
-----
-----
```

Lasso Regression Training Accuracy

```
MSE: 0.43
MAE: 0.52
RMSE: 0.65
R2 Score: 0.57
Adjusted R2: 0.56
-----
```

Lasso Regression Test Accuracy

```
MSE: 0.43
MAE: 0.51
RMSE: 0.65
R2 Score: 0.58
Adjusted R2: 0.55
-----
-----
```

```
# Assuming y_pred_train, y_pred_train_ridge, and y_pred_train_lasso
are your prediction arrays
```

```
actual_values = y_train.values.reshape((-1,))
```

```
predicted_values = [
    y_pred_train.reshape((-1,)),
    y_pred_train_ridge.reshape((-1,)),
    y_pred_train_lasso.reshape((-1,)) # Change from reshape(()) to
reshape((-1,))
]
```

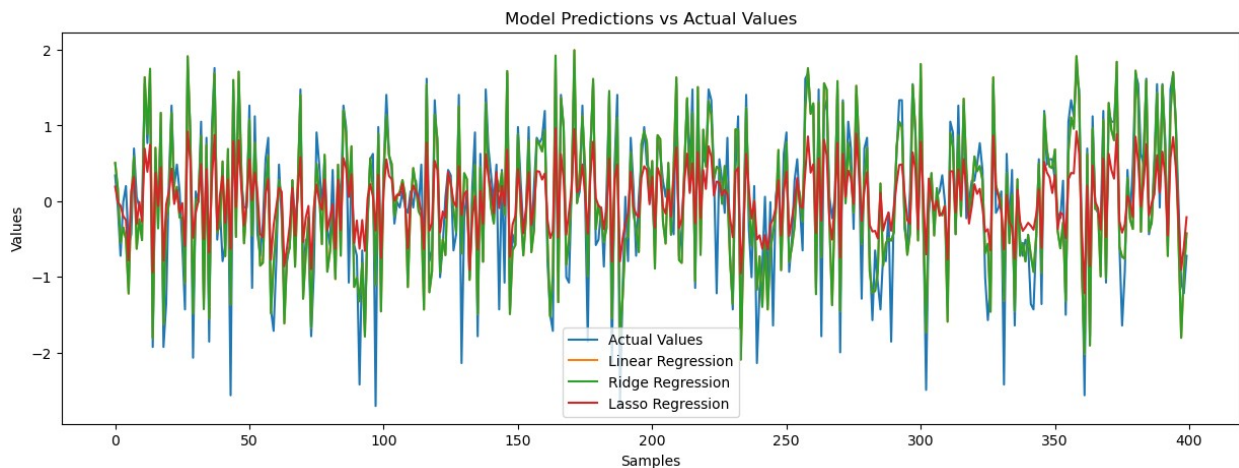
```
# Continue with your plotting or analysis
```

```
model_names = ['Linear Regression', 'Ridge Regression', 'Lasso
Regression']
```

```
plt.figure(figsize=(15, 5))

# Example plotting (assuming you want to plot these values)
plt.plot(actual_values, label='Actual Values')
for i, preds in enumerate(predicted_values):
    plt.plot(preds, label=model_names[i])

plt.legend()
plt.title('Model Predictions vs Actual Values')
plt.xlabel('Samples')
plt.ylabel('Values')
plt.show()
```



Elastic-Net Regression

```
ElasticNet_model = ElasticNet(alpha=0.108)
ElasticNet_model.fit(x_train , y_train)

ElasticNet(alpha=0.108)

y_pred_train_el = ElasticNet_model.predict(x_train)
y_pred_test_el = ElasticNet_model.predict(x_test)

train_R2 = ElasticNet_model.score(x_train,y_train)
test_R2 = ElasticNet_model.score(x_test,y_test)
train_R2 , test_R2

(0.814348667393518, 0.8153699952125263)

train_R2 = ElasticNet_model.score(x_train,y_train)
test_R2 = ElasticNet_model.score(x_test,y_test)
train_R2 , test_R2

(0.814348667393518, 0.8153699952125263)
```

```

en_model_weights = pd.DataFrame(ElasticNet_model.coef_.reshape(1, -
1), columns=jb.columns[:-1])
en_model_weights["Intercept"] = ElasticNet_model.intercept_
en_model_weights

```

	GRE Score	TOEFL Score	University Rating	SOP	LOR
CGPA \					
0	0.194912	0.13349	0.025508	0.02075	0.091694
	0.418398				

	Research	Intercept
0	0.058465	0.007728

```

print('ElasticNet Regression Training Accuracy\n')
model_evaluation(y_train.values, y_pred_train_el, ElasticNet_model)
print('*'*25)
print('\nElasticNet Regression Test Accuracy\n')
model_evaluation(y_test.values, y_pred_test_el, ElasticNet_model)
print('---'*25)

```

ElasticNet Regression Training Accuracy

```

MSE: 0.18
MAE: 0.31
RMSE: 0.43
R2 Score: 0.81
Adjusted R2: 0.81
*****

```

ElasticNet Regression Test Accuracy

```

MSE: 0.19
MAE: 0.3
RMSE: 0.44
R2 Score: 0.82
Adjusted R2: 0.81
-----
-----

```

```

model_major_weights = {"Linear Model":lr_model_weights,
"Ridge Model":ridge_model_weights,
"Lasso Model":lasso_model_weights,
"Elastic_Net":en_model_weights}

```

excluding w0-intercept

```

plt.figure(figsize=(25, 5))
i = 1

```

```

for model, data in model_major_weights.items():
    model_weights_data = data.melt() # Melt the DataFrame for

```


```

plotting
plt.subplot(1, 4, i) # Create subplots
sns.barplot(
    data=model_weights_data[:-1].sort_values(by='value',
ascending=False), # Sort values
    y='variable',
    x='value',
    width=0.4,
    palette=['darkgreen', 'g', 'green', 'limegreen', 'seagreen',
'mediumseagreen']
)

plt.xlabel('Features')
plt.ylabel('Coefficient/Weight')
plt.title(f'{model} Coefficients / Weights')
i += 1

sns.despine() # Remove top and right spines from plots
plt.show() # Display the plots

```



Regression Analysis Summary:

1. By conducting regression analysis, it's evident that CGPA emerges as the most influential feature in predicting admission chances.
2. Additionally, GRE and TOEFL scores also holds significant importance.
3. Following the initial regression model, a thorough check for multicollinearity was performed, revealing VIF scores consistently below 5, indicative of low multicollinearity among predictors.
4. Despite the absence of high multicollinearity, it's noteworthy that the residuals do not conform perfectly to a normal distribution. Furthermore, the residual plots indicate some level of heteroscedasticity.
5. After exploring involving regularized models such as Ridge and Lasso regression showed comparable results to the Linear Regression Model.
6. Moreover, employing ElasticNet (L1+L2) regression yielded results consistent with the other regression models.

Recommendation

7. Encourage students to focus on improving GRE scores, CGPA, and Letters of Recommendation (LOR), as these factors influence a lot your chances of admission.
8. Beyond academic metrics applicants can also add like extracurricular achievements, personal statements, and diversity factors.
9. We can enhance our predictive model by adding other important and diverse features like Work-experiece, internships or extra- curriculum activites.

